BrandGuard

An intelligent brand-governance platform that reviews content before it is published, explains every issue and suggests safe corrections.

Round 1 Submission
Problem Statement: Brand Governance (Inkloom AI)

Team Name	**THE_OGS**
Team Members	1. **Trusha (Team Leader)** 2. **Manvi Chaudhary** 3. **Rohit Bansal** 4. **Tanmay Sethi**

Contents and Requirement Map

Every item required by the submission instructions is covered, in the order requested. The right-hand column shows where the judging criteria are answered.

#	Required section	Page area	Judging criterion served
1	Understanding of the problem statement	Section 1	Understanding of the problem (25)
2	Users affected by the problem	Section 2	Understanding of the problem (25)
3	Proposed solution (includes originality)	Section 3	Solution (25), Originality (15)
4	Complete product workflow	Section 4	Solution (25)
5	Main features to build	Section 5	Solution (25), Feasibility (15)
6	Technology stack and technical plan	Section 6	Technical plan (20)
7	Responsibilities of each team member	Section 7	Feasibility (15)
8	Development plan across three rounds (includes feasibility)	Section 8	Feasibility (15)
9	Expected impact	Section 9	Solution (25)
10	Future scope after the event	Section 10	Originality (15)
11	Complete team and participant details	Section 11	Submission requirement

Ship it clean. Your brand, your voice.
Summary in one line: BrandGuard checks content against approved guidelines, facts and assets before it goes live, explains every issue, suggests fixes that keep the creator’s message, and leaves the final decision to a human reviewer.

The team has designed the rule catalogue, the approved-fact structure and the review workflow on paper and chosen a plan that is realistic for the build window. All development will happen during the official build period.

1. Understanding of the Problem Statement

What the problem statement asks. Growing organisations publish content through many employees, agencies and AI tools. Because many hands touch the content, it is hard to keep a consistent and trustworthy brand. The brief asks for a platform that reviews content before publication, checks it against approved guidelines, facts and assets, detects specific categories of issue, explains each issue, suggests corrections, supports a collaborative approval workflow, works across formats, and gives insights, all while balancing automation with human control and protecting confidential material.

The root problem. The issue is not that people are careless. Brand knowledge lives in PDFs, slide decks and people's heads, so nobody can check everything at the speed content is now produced, especially with AI tools generating drafts in seconds. Mistakes are found after they are public.

How the problem hurts

Impact	What happens
Lost credibility	A wrong claim or outdated figure makes audiences doubt everything else the brand says
Legal and compliance exposure	Unsupported promotional claims ("guaranteed", "best", "#1") can create regulatory risk
Inconsistent identity	Different voices, terms and visuals make the brand look disorganised
Excluded audiences	Missing alt text, low contrast and complex language shut out readers
Information leaks	Codenames or private data can slip into public posts
Wasted effort	Reviewers repeat the same corrections; creators feel blocked by unexplained rejections

The part we intend to solve

We will solve the **complete pre-publication review loop** for text-based content (social posts, emails, web copy) and visual specifications: detect, explain, fix, approve, learn. We deliberately do **not** try to build a full content-management system, a design tool, or a fully automatic publisher. We focus on depth and trustworthiness of the review itself.

2. Users Affected by the Problem

User	Their pain today	What BrandGuard gives them
Content creator (employee, freelancer, agency writer)	Unsure what is allowed; receives vague rejections late; rewrites repeatedly	Instant, explained feedback with one-click fixes that keep their message
Marketing / brand manager	Spends time policing basic errors; cannot see patterns across agencies	Routine issues pre-solved; dashboard of recurring mistakes by team and agency
Legal / compliance reviewer	Sees content too late and too often; must hunt for risky claims	Only risky items escalate, each with evidence, rule and audit trail
Brand administrator	Guidelines are ignored or misread; no feedback on which rules confuse people	Guideline-health report shows which rules are unclear and need examples
External agency	Receives guidelines as a PDF; is corrected after delivery	Checks drafts against the same rules before sending work to the client
Audience and customers	Receive wrong, inconsistent or inaccessible content	More accurate, accessible and consistent communication

Primary persona

Priya, content creator at a growing energy company. She drafts social posts and emails daily, sometimes with an AI tool. She wants to publish fast but cannot memorise 40 pages of guidelines. She needs feedback that tells her what is wrong, why, and how to fix it without losing her message.

3. Proposed Solution

BrandGuard is a review platform with three layers that work together: a **rule and fact engine** that detects issues reliably, an **explanation and suggestion layer** that tells creators what to do, and a **workflow and insight layer** that keeps humans in control and improves the guidelines over time.

Issue categories detected

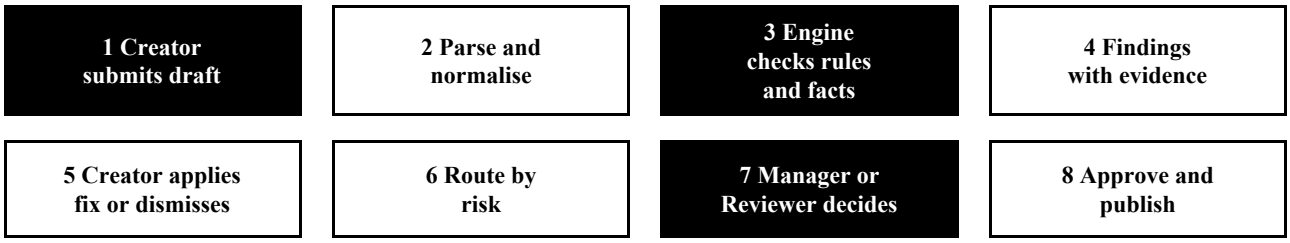
Category	Examples
Tone	Hype words, shouting in capitals, repeated exclamation marks
Terminology	Banned versus preferred terms (for example "users" versus "customers")
Messaging	Off-message positioning, missing mandatory statements
Unsupported claims	"Best", "guaranteed", "#1" with no approved evidence
Outdated or wrong facts	Warranty years, savings percentages, company facts that differ from the approved fact library
Visual consistency	Off-palette colours, unapproved fonts, logo misuse
Accessibility	Missing alt text, low colour contrast, vague link text, very long sentences, unreadable hashtags
Confidentiality	Project codenames, access keys, personal email addresses, "internal only" text

What makes our approach different (originality)

- **Fix, do not just flag.** Every finding comes with a correction that preserves the creator's original purpose. Example: "save up to 50%" becomes "typically save 20-30%", still a savings message, now truthful.
- **Verify-after-fix loop.** Any suggested rewrite, including AI-generated ones, is re-run through the same rule engine before the creator sees it, so a suggestion can never introduce a new violation.
- **Facts come from a governed library, not from the AI.** The AI may only use approved facts, each with a source and a review date. It cannot invent numbers.
- **Confidence-based routing.** Each finding carries a confidence level. Safe fixes can be applied in one click; uncertain or high-risk items go to a human reviewer.
- **The platform audits the guidelines too.** Rules that people dismiss often are flagged as probably unclear, so the guideline is fixed instead of blaming creators.
- **Dismissals need a reason.** This turns disagreement into data and keeps accountability.

4. Complete Product Workflow

End-to-end flow



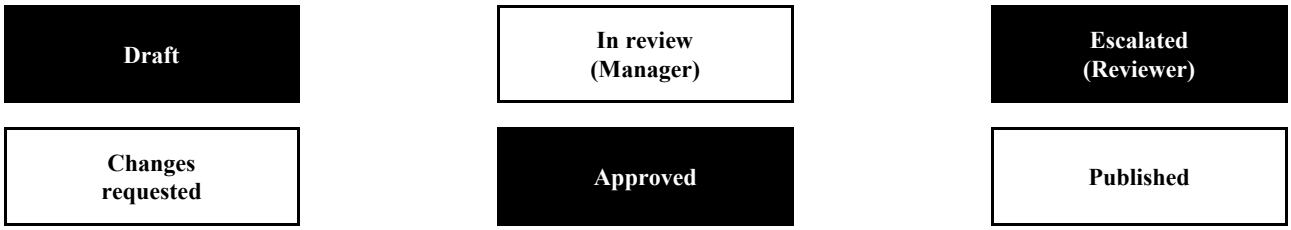
Step by step

- **Submit.** The creator pastes text, uploads a file or sends content from a connected tool, and selects the format and channel.
- **Check.** The engine compares the content with the rules, approved fact library, palette, fonts and watch-lists.
- **Explain.** Each finding shows the exact text found, the rule, why it matters, a suggested fix and a confidence level.
- **Fix.** The creator applies a fix with one click, edits it, or dismisses it with a stated reason. The content is re-checked after every change.
- **Route.** The system decides who must approve, based on severity (table below).
- **Decide.** The approver can approve, request changes with comments, or (for critical issues) override only with written justification.
- **Publish and learn.** Approved content is marked ready for publishing. All actions feed the audit trail and the insight dashboard.

Human-in-control routing

Situation	Who decides
No open issues, or only minor findings	Manager can approve
Major issue: claims, facts, visuals, accessibility failure	Authorised Reviewer must approve
Critical issue: confidential information	Reviewer only; override needs written justification
Creator disagrees with a finding	Dismisses with a reason; logged and counted in insights
Any status change	Recorded with who, when and why

Approval states



5. Main Features to Build

Feature	What it does	Priority
Review engine	Detects tone, terminology, messaging, claim, fact, visual, accessibility and confidentiality issues	Core
Explainable findings	Shows evidence, rule ID, reason, suggestion and confidence for every issue	Core
Intent-preserving fixes	One-click corrections that keep the creator's meaning; re-checked after applying	Core
Approved fact library	Facts with sources and review dates used for claim and figure checking	Core
Approval workflow	Creator, Manager and Reviewer roles; request changes; versions; comments	Core
Audit trail	Record of every review, edit, approval, dismissal and override	Core
Confidentiality protection	Watch-list scan, role-based access, local or in-organisation processing	Core
Insights dashboard	Repeated mistakes by rule, team and content type; unclear-guideline detector	Core
Guideline manager	Admin interface to switch rules on or off, edit terms and facts	Core
Grounded AI suggestions	Language-model rewrites limited to approved facts, with verify-after-fix loop	Round 3
Multi-format ingestion	Images (alt text, palette), PDFs and decks in addition to text	Round 3 (stretch)
Slack and CMS connectors	Check drafts where teams already work	Future scope

Example finding

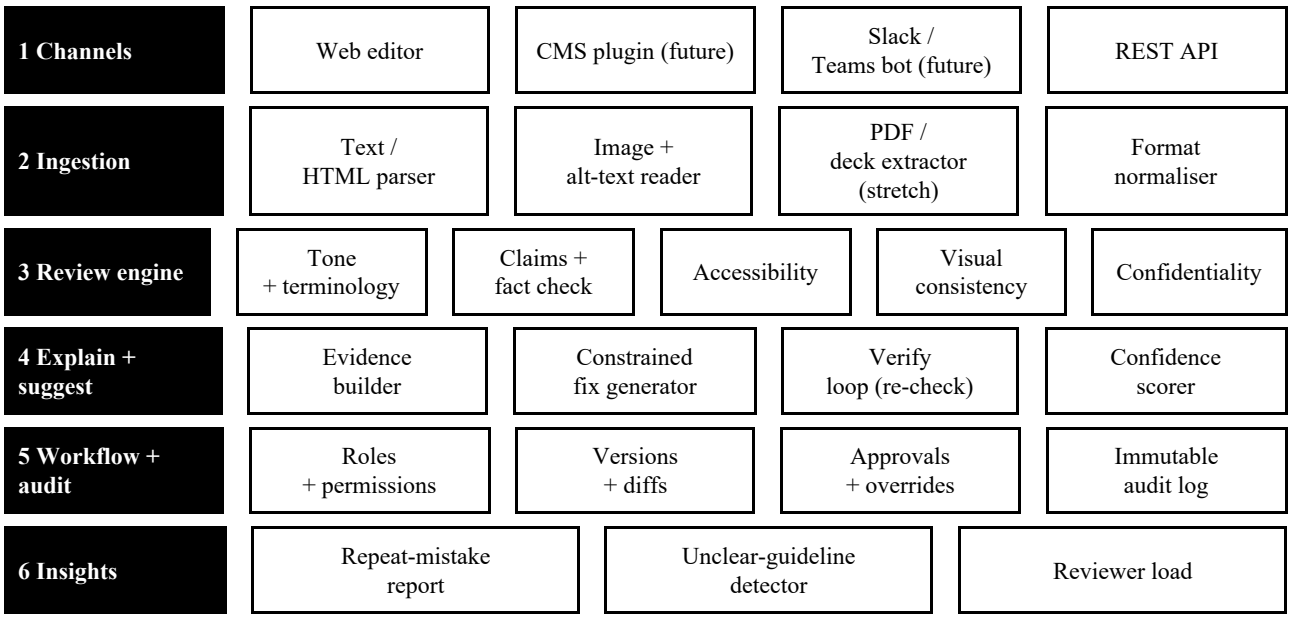
Draft: "Our amazing panels are the BEST on the market! Save up to 50% on your bills, guaranteed."

	Finding 1: Claims (major)	Finding 2: Facts (major)
Evidence found	"guaranteed"	"save up to 50%"
Why it matters	No approved evidence supports a guarantee	Approved evidence supports 20-30%, not 50%
Suggested fix	Refer to what is actually guaranteed, such as the 25-year performance warranty	"Typically save 20-30%" (one-click apply)
Intent preserved?	Yes, still confident and benefit-led	Yes, still a savings message, now accurate
Illustrative values	80%, human check recommended	95%, safe to apply

6. Technology Stack and Technical Plan

Layer	Choice	Why we chose it
Review engine	Python, rule-based checks first; spaCy and scikit-learn where useful	Rules are deterministic and explainable, which judges and reviewers can trust; ML is added only where rules are not enough
AI layer	Large language model through an API, behind a retrieval layer	Good at rewriting while keeping tone; restricted to approved facts so it cannot invent information
Backend	FastAPI	Fast to build, typed, automatic API documentation
Database	PostgreSQL	Reliable relational store for content versions, rules, facts and the audit log
Frontend	Next.js / React with accessible components	Rich review interface; accessible by default, which matches the product's purpose
Workflow	State machine in the backend with role-based access	Clear, testable approval rules; no extra service to run
Deployment	Vercel , Render	Free hosting that gives judges a public link, with no installation needed.

System architecture



A **Brand Knowledge Base** (guidelines, fact library with sources and expiry dates, palette, fonts, logos, message library) feeds the engine. Only Brand Admins can change it, and every change is versioned.

6. Technology Stack and Technical Plan (continued)

Most technically challenging part

Suggesting corrections that preserve the creator's intent without introducing new errors or invented facts. Detecting a bad claim is straightforward; rewriting it so it is truthful, on-brand and still says what the creator meant is hard, and a careless AI rewrite can add new unsupported claims.

How our architecture handles it

- **Two-tier fixes.** Simple, high-confidence issues (a wrong number, a banned term) use deterministic replacements taken from the fact library and term dictionary. No AI is involved, so the result is predictable.
- **Constrained AI for complex rewrites.** For tone and messaging, the language model receives only the original text, the violated rule and the approved facts it is allowed to use.
- **Verify-after-fix loop.** Every suggestion is re-run through the review engine. If it creates a new violation, it is discarded and the creator sees a plain explanation instead of a risky rewrite.
- **Confidence and escalation.** Low-confidence suggestions are labelled and routed to a human reviewer.
- **Measured quality.** A labelled test set of good and bad content (with known seeded issues) is used to track how many issues are detected and how many false alarms occur.

Other technical risks

Risk	Mitigation
False positives annoy creators	Confidence scores, dismiss-with-reason, and rule tuning from the insight dashboard
Facts become outdated	Each fact has a source and review date; admins get expiry alerts
Confidential content exposure	In-organisation processing, role-based access, no training on customer content, redaction of flagged items
AI invents information	Retrieval limited to approved facts plus the verify-after-fix loop

7. Responsibilities of Each Team Member

Each member owns a clear area and a measurable deliverable so that individual contributions can be assessed. Replace bracketed names with the real team.

Member	Role	Owns	Round deliverables
Trusha (Team Leader)	Product lead and integration	Scope, project charter, architecture decisions, integration of all modules, final presentation and demo	R1: this document. R2: integrated build. R3: final demo and pitch
Rohit Bansal	Rule and AI engine	Rule catalogue, fact checks, fix generation, verify-after-fix loop, AI guardrails, test set	R2: rule engine with explanations. R3: grounded AI suggestions
Tanmay Sethi	Backend and workflow	Database, API, roles, approvals, audit trail, insights calculations, deployment	R2: API and database. R3: workflow, insights, deployment
Manvi Chaudhary	Frontend and experience	Review interface, accessibility of the product itself, dashboards, user testing, documentation	R2: review screen. R3: workflow screens, insights, polish

Team practices

- One shared GitHub repository with issues assigned per member and pull-request reviews.
- Short weekly checkpoint led by the Team Leader against the plan in Section 8.
- Each member records their own commits and notes to support individual contribution evaluation.

8. Development Plan Across All Three Rounds

Round	Goal	What will be done	Done when
Round 1	Plan and prove feasibility	Problem analysis, user research, solution design, architecture, planned rule catalogue and fact-library structure	Judges can understand the problem, solution and plan
Round 2	Working core	Python/FastAPI engine with full rule set; PostgreSQL storage; explainable findings; one-click deterministic fixes; review interface; labelled test set	A user can submit content and receive correct, explained findings and fixes
Round 3	Complete and polish	Approval workflow with roles and audit trail; insights dashboard and unclear-guideline detector; grounded AI suggestions with verify loop; deployed link; README; demo script	End-to-end demo from submission to approval, plus insights, runs from a public deployed link

Feasibility and scope

- **Focused scope.** We limit the build to text content and visual specifications. Images, PDFs and decks are a stretch goal, and Slack/CMS connectors are future scope.
- **Rules before AI.** The deterministic engine is useful on its own. If time runs short, the AI layer can be reduced without breaking the product.
- **Clearly scoped.** The core engine, workflow and insights already work in a proof of concept, so the team is extending a working base.
- **Standard tools.** Every technology in the stack is well documented and familiar to student teams.
- **Prioritised backlog.** Core features come first, stretch items only after the core demo is stable.

Risks and fallback

Risk	Fallback
AI integration takes longer than planned	Ship with deterministic fixes and templated explanations
Image or PDF parsing proves too complex	Accept pasted text and a visual-spec form instead
A team member is unavailable	Cross-training on modules and documented interfaces

Three improvements we commit to delivering

- Grounded AI rewrites with a verify-after-fix loop
- Unclear-guideline detector based on dismissal patterns
- Image checking (alt text and colour palette)

9. Expected Impact

The table below describe the outcomes we expect and how we will measure them. They are goals to validate in testing and pilots, not results we have already measured.

Who benefits	Expected outcome	How we will measure it
Brand and marketing teams	Fewer errors reach the public; less time spent on routine corrections	Issues caught before publication versus a seeded test set; reviewer time per item
Creators and agencies	Faster, clearer feedback; fewer rounds of rework	Share of findings accepted or fixed; number of resubmissions per item
Legal and compliance	Attention focused on genuinely risky items	Escalation rate; share of escalations that were truly high risk
Brand administrators	Clearer guidelines over time	Dismissal rate per rule falling after rules are improved
Audiences	More accurate, consistent and accessible content	Accessibility issues remaining at publication; factual errors remaining at publication
The organisation	Stronger trust, lower legal and reputational risk, fewer leaks	Confidential items blocked in test set; audit trail completeness

10. Future Scope After the Event

- **More formats:** image analysis (logo use, colour, text in images), PDFs, slide decks and video captions.
- **Integrations:** Slack and Microsoft Teams bots, CMS plugins, and checks inside design and email tools.
- **Multi-language brand voice:** localised terminology and tone profiles.
- **Learning from outcomes:** using accepted and dismissed findings to refine rules and personalise feedback, always with human approval of rule changes.
- **Agency portal:** a controlled space where external partners self-check work before delivery.
- **Enterprise readiness:** single sign-on, fine-grained permissions, data residency controls, and compliance reports.
- **Pre-flight for AI tools:** a gateway that checks content produced by AI assistants before it leaves the organisation.

1. Complete Team and Participant Details

Field	Details
Team name (exactly as in form)	THE_OGS
Project title	BrandGuard: Intelligent Brand Governance Platform
Problem statement	Brand Governance (Inkloom AI)
Round	Round 1

Role	Full name (as on certificate)	Email	Phone	College / organisation
Team Leader	Trusha	trushaa15@gmail.com	7652975339	UIET Chandigarh
Member 2	Manvi Chaudhary	manvi.chaudhary052@gmail.com	7696679374	UIET Chandigarh
Member 3	Rohit Bansal	coderrohit02@gmail.com	8360690282	UIET Chandigarh
Member 4	Tanmay Sethi	tanmaysethi1221@gmail.com	9306688934	UIET Chandigarh