BrandPilot Intelligence

AI Marketing Department | Project Blueprint

One platform to coordinate content, brand governance, audience engagement and marketing performance.

Product vision
BrandPilot helps organisations handle more marketing work with fewer repetitive manual tasks and fewer disconnected tools. It brings specialised AI workflows into one shared system, with approved brand information, clear permissions and human review where it matters.

Operating principle
Automate repeatable tasks, assist with judgement-heavy work, and require authorised human approval for legal, sensitive or high-impact decisions.

Plan → Create → Verify → Approve → Publish → Engage → Measure → Improve

1. The Problem

Growing businesses often depend on a mix of social media managers, content creators, copywriters, designers, analysts, community managers, sales staff and compliance reviewers. Work is spread across chats, spreadsheets, design tools, schedulers and analytics dashboards.

The result is duplicated effort, delayed approvals, inconsistent messaging, outdated offers, missed customer enquiries and limited visibility into which marketing activities produce results.

BrandPilot's goal is not to claim that every employee can be removed. It is to consolidate repeatable work, reduce tool switching and allow a smaller team to coordinate more campaigns with clear accountability.

2. Product Positioning

Positioning statement: BrandPilot is an AI-powered marketing operations platform that helps businesses plan, create, validate, publish and measure marketing work from one governed workspace.

Core promise: Reduce repetitive marketing workload and dependence on disconnected tools while keeping brand facts, approvals, customer communications and results traceable.

What makes it different

- Shared brand memory: every module uses the same approved facts, offers, voice, visual rules and assets.
- Connected workflows: a change to an offer can trigger checks across drafts, scheduled posts and response templates.
- Evidence-linked review: warnings identify the claim, rule, source and required correction.
- Controlled automation: exact rules govern when the platform can publish or respond; uncertain cases go to a person.
- Outcome tracking: the platform measures time saved, review quality, engagement and qualified conversions rather than promising virality.

3. Roles BrandPilot Can Support

Role	Tasks the platform can automate or assist with	Human responsibility that remains
Content creator / copywriter	Hooks, scripts, captions, carousels, story sequences, CTAs, keyword ideas and format variations.	Original perspective, creative direction and final quality judgement.
Social media manager	Content calendar, briefs, reminders, scheduling recommendations and campaign status.	Priorities, resource allocation and strategy.
Graphic designer / brand manager	Brand-safe template variations, asset retrieval, logo/colour/type checks and design briefs.	Art direction, complex design and subjective decisions.
Community manager	FAQ replies, comment categorisation, approved keyword DMs and escalation routing.	Sensitive conversations, conflict resolution and relationship-building.
Marketing analyst	Baseline comparisons, campaign summaries, anomaly alerts and experiment suggestions.	Interpreting business context and selecting consequential actions.
Legal / compliance reviewer	Flag unsupported claims, expired offers, missing disclosures and restricted wording.	Legal interpretation, approval and accountability.
Sales / lead qualification	Intent detection, approved product answers, lead routing and next-step prompts.	Negotiation, complex objections and closing strategic deals.
Marketing coordinator	Assignments, deadlines, approval reminders, status reports and handoffs.	Resolving conflicts and coordinating unusual cases.

The correct promise is role consolidation and workload reduction, not guaranteed elimination of jobs. The amount of work saved depends on the business, channel access, quality requirements and human approval policy.

4. Three Levels of Work

Mode	Meaning	Example
Automate	The system performs a predictable task under a defined rule.	Check an offer expiry date or send a pre-approved reply to an allowed keyword trigger.
Assist	AI proposes work for a person to inspect or edit.	Generate three hooks or explain why a post underperformed.
Human-controlled	A person must approve the action or decide the exception.	Approve a regulated claim, major budget change or sensitive complaint response.

5. One Shared Brand Knowledge Base

All modules must read from a version-controlled source of truth. Do not let individual AI agents invent or maintain conflicting versions of important facts.

Record	Required information
Brand identity	Positioning, audience, tone attributes, preferred and prohibited terms, approved examples.
Visual system	Logo files and variants, colours, fonts, templates, spacing rules and asset usage rights.
Product facts	Features, limitations, supported use cases, pricing, availability, warranty/refund details where applicable.
Offers	Price, discount, start/end time, eligibility, exclusions, stock conditions, owner and approval status.
Claims and proof	Exact claim, source, evidence, permitted wording, qualifications, owner and expiry/review date.
Audience rules	Relevant segments, problem, intent signals, exclusions and permitted personalisation.
Platform rules	Connected account, API capabilities, approved triggers, publishing permissions and rate limits.

Every important record needs an owner, version, approval status and effective date. When a record changes, dependent drafts and templates should be rechecked.

6. Content Planning and Production

The system should not generate a generic daily schedule. It should plan from the campaign objective, audience, channel, available creative assets, offer status, team capacity and past performance.

Output	Required fields
Calendar item	Date/time, channel, format, objective, audience, funnel stage, owner and approval deadline.
Reel	Opening hook, scene-by-scene script, on-screen text, proof, duration target, cover and CTA.
Carousel / post	Slide sequence or headline, supporting facts, source links, visual brief and CTA.
Story sequence	Frame order, poll/question, reply route, link, offer terms and final action.
Caption package	Short/long versions, relevant keywords, platform-appropriate hashtags when useful and accessibility text.
Review record	Source-backed claims, brand findings, accessibility findings, reviewer and final approved version.

When to publish

Use the account's own historical performance by format, day, time zone and objective. If there is insufficient data, label the recommendation as a test. Do not present a universal posting time as a guarantee.

Hooks and CTAs

Generate several options tied to the audience's real problem and the content's actual payoff. A hook must not promise a result the content cannot substantiate. Each post should have one clear next action, and the destination or promised asset must be tested before publication.

7. What Is Going Viral, and What Is Not?

Virality cannot be predicted with certainty. BrandPilot should compare content against the account's own baseline and explain which metrics support its classification.

- Video: watch time, completion and retention data where available.
- Discovery: reach, impressions and non-follower reach where supplied by the platform.
- Intent: shares, saves, profile visits, meaningful replies, link clicks and qualified leads.
- Business outcome: sign-ups, purchases or contribution margin where tracking is reliable.
- Context: topic, hook, length, format, audience, CTA, paid promotion and campaign timing.

Result label	Decision rule
Above baseline	A chosen metric exceeds a defined comparison baseline with sufficient observations.
Below baseline	A chosen metric falls below the baseline; report possible factors without claiming causation.
Inconclusive	Insufficient data, incomplete analytics or too few meaningful outcomes.
Potential pattern	Similar results appear across multiple comparable posts; show sample size and exceptions.

The output should propose one test at a time. For example, test a different opening while keeping the topic and CTA similar. Do not claim that a trend or format guarantees reach.

8. Brand Visuals and Identity

- Check logo variant, clear space, minimum size, approved background and prohibited alterations.
- Validate exact brand colour values, typography rules, canvas dimensions and text overflow where technically measurable.
- Use approved templates and current assets; retain asset source, licence, permitted usage and expiry.
- Flag possible tone or visual-style mismatches as review suggestions when they require subjective judgement.
- Never describe an AI visual as approved simply because it resembles the brand.

9. Factual, Offer and Legal Review

The release gate checks drafts against approved facts and routes defined risk categories to authorised reviewers. It is a risk-reduction workflow, not a guarantee of legal compliance.

Risk	System behaviour
Unsupported performance claim	Require evidence or a rewrite before release.
Outdated price or product detail	Flag the exact text and point to the current approved record.
Expired promotion	Block scheduling or publishing under the configured policy.
Missing eligibility or offer terms	Identify the missing condition and request correction.
Testimonial or comparative claim	Check evidence and disclosures; require review when policy says so.
Regulated or sensitive content	Route to an authorised human reviewer; do not auto-approve.
Unclear evidence	Mark “not verified”; never fabricate a source or citation.

Each finding should show the exact wording, severity, rule, source, suggested correction, reviewer and resolution status. Editing an approved version must trigger the relevant checks again.

10. Comments, DMs and ManyChat-Style Automation

The engagement module uses exact triggers or supported intent classification with approved response templates. Actual functionality depends on the connected platform's current API permissions and policies.

Trigger	Action	Control
Comment keyword such as CHECKLIST	Send the approved resource link by DM where supported.	Verify the asset and log delivery result.
Price question	Return current approved price and terms.	Read the live approved record, not an old prompt.
Demo request	Ask a minimal qualifying question or share the approved booking link.	Never invent availability.
Story reply	Use story context to select an approved response or ask for clarification.	Escalate low-confidence cases.
Complaint / refund	Acknowledge and route to support.	No invented policy or automatic denial.
Legal or sensitive issue	Pause automated handling and notify the responsible person.	Human review required.

- Include opt-out handling, per-user frequency limits, duplicate-event protection and a global pause switch.
- Use test mode before activating rules; record trigger, template version, delivery status and handoff.
- Do not claim a message was delivered until the connected service confirms it.
- Do not use automation to spam, evade platform limits or contact people without the required permission.

11. Audience, Leads and Offers

Audience selection

Use relevant business criteria such as stated problem, role, industry, company size when appropriate, geography where necessary and explicit intent. Avoid sensitive-trait inference or treating weak behavioural signals as certain facts.

Lead qualification

- Need: does the stated problem match a supported use case?
- Fit: can the product genuinely serve the requested use case?
- Intent: did the person ask for pricing, a demo, availability or next steps?
- Routing: send qualified leads to the right owner; send existing customer problems to support.
- Uncertainty: ask one useful clarifying question rather than pretending to know.

Offer timing

- Require a defined objective, owner, audience, terms, duration and approved financial guardrails.
- Do not invent scarcity, deadlines, testimonials, discount eligibility or expected results.
- Make terms consistent across posts, DMs, landing pages and checkout.
- Measure qualified conversions and margin where available, not clicks alone.

12. Workflow and Human Approval

Stage	Required result
Brief	Objective, audience, channel, offer and evidence needs defined.
Draft	Content and assets created from approved brand information.
Automated checks	Specific findings with evidence and suggested corrections.
Revision	Author fixes findings or records a permitted exception.
Human approval	Required manager or compliance approvals recorded for that version.
Publish / schedule	Connected account and final version confirmed.
Monitor	Results, delivery errors and audience responses collected where available.
Improve	Findings feed future briefs, tests and guideline clarification.

An author must not bypass required separation of duties. Emergency exceptions must identify the approver, reason, scope and expiry.

13. System Architecture

Module	Purpose
Workspace and permissions	Organisation, client isolation, roles and account access.
Brand knowledge base	Versioned facts, guidelines, assets, offers and evidence.
Planning and content studio	Calendar, briefs, hooks, scripts, captions and creative instructions.
Rules engine	Exact checks for dates, prices, terms, required disclosures and permissions.
AI review layer	Tone and meaning analysis, grounded suggestions and explanations.
Approval and audit	Versioning, review queue, decision records and history.
Engagement inbox	Comment/DM classification, templates, keyword rules and human handoffs.
Channel connectors	Authorised publishing, message events and analytics imports.
Measurement service	Baselines, experiment tracking, campaign reporting and outcome records.
Notifications	Approval reminders, expiring offers, integration failures and escalation alerts.

Use deterministic code for exact rules and permissions. Use AI for language interpretation and drafting. Store the evidence and rule version behind each important warning.

14. Confidentiality, Security and Reliability

- Enforce client workspace isolation at both API and database layers.
- Use role-based access and limit exports of confidential content and customer messages.
- Store integration tokens in a secrets manager, not source code or prompts.
- Encrypt data in transit and at rest; restrict backup access.
- Record sensitive access, changes, approvals, automated replies and publishing actions.
- Document which AI providers process content and obtain organisational approval before sending confidential materials.
- Set retention and deletion rules for drafts, messages, analytics, assets and audit records.
- Retry transient failures safely, prevent duplicate posts/replies and alert on expired credentials.
- Provide authorised users with a way to pause publishing, revoke connections and disable automation.

15. Technology and Integration Boundaries

Layer	Initial choice / rule
Frontend	React or Next.js for calendar, editor, approvals, inbox and reports.
Backend	Choose one primary backend, such as FastAPI or Node.js; avoid duplicated business logic.
Database	PostgreSQL for users, brand facts, offers, versions, findings and approvals.
Asset storage	Private object storage with restricted and time-limited access.
Background jobs	Queue and workers for scheduled checks, publishing jobs and analytics imports.
AI	LLM with retrieval from approved records and explicit output validation.
Social channels	Official APIs or approved providers; verify each specific permission before promising a feature.
Monitoring	Structured logs, metrics, alerts and safe retry handling.

Reading analytics does not automatically grant permission to publish, reply to comments or send DMs. Verify each capability for the exact channel and account type.

16. How to Prove Time and Cost Savings

The product should measure workload reduction instead of claiming that a fixed number of jobs can always be removed.

Measure	How to calculate
Time saved per task	Baseline median minutes minus BrandPilot-assisted median minutes for the same task type.
Approval turnaround	Time from submission to final approval, compared before and after adoption.
Correction rate	Share of published items requiring preventable correction, using a defined classification.
Automation reliability	Successful eligible actions divided by eligible attempted actions, with failures and duplicates tracked.
Cost comparison	Software and operating cost plus remaining human labour, compared with the prior workflow.
Business outcome	Qualified leads, conversions and margin against an appropriate baseline; avoid attributing all change to the tool.

A customer-specific business case should compare current staff/tool costs with the actual subscription, integration, review and oversight costs after adoption. Savings vary by workload and should be demonstrated in a pilot.

17. Data Model: First Draft

Entity	Key data
Organisation / Workspace	IDs, settings, retention policy and access boundaries.
User / Role / Permission	Identity, workspace, assigned role and permissions.
BrandGuideline / BrandAsset	Version, rules, file reference, rights, expiry and approval.
ApprovedFact / ClaimEvidence	Claim, value, source, owner, validity dates and status.
Offer / OfferTerm	Price, discount, eligibility, exclusions, start/end and approval.
Campaign / AudienceSegment	Goal, segment criteria, budget guardrails and owner.
ContentItem / ContentVersion	Channel, format, copy, assets, source links, version and status.
ValidationFinding / Approval	Rule, severity, evidence, reviewer, decision and timestamp.
PublishingJob / ChannelAccount	Account, scheduled time, external ID, result and error.
EngagementEvent / ReplyRule	Trigger, intent, template version, action result and escalation.
MetricSnapshot / Experiment	Metric, value, time window, baseline, sample size and test.
AuditEvent	Actor, action, target, timestamp, result and correlation ID.

18. Development Plan

Phase	Deliverable	Exit test
0. Scope	First customer segment, problem, pilot goal and boundaries.	A specific buyer and measurable pain are agreed.
1. Requirements	User stories, workflow and acceptance tests.	Every core feature has a test.
2. Brand/fact base	Versioned facts, offers, assets and rules.	Expired and outdated records are caught.
3. Governance MVP	Draft intake, findings, revisions and approvals.	A user can trace each warning to evidence.
4. Content planning	Calendar, briefs and content package generation.	Drafts use approved facts and remain editable.
5. One-channel integration	One verified publishing or engagement flow.	End-to-end test passes without duplicate actions.
6. Automation and handoff	Keyword rules, limits, pause switch and human queue.	Unknown and sensitive cases escalate.
7. Analytics	Baselines, experiment reports and outcome metrics.	Results reconcile with source data.
8. Security and pilot	Isolation, retention, monitoring and customer testing.	Security tests pass and savings are measured.

Recommended MVP: one customer segment, one social channel, shared brand/fact records, content calendar, pre-publication checks, approval workflow and a small number of approved engagement triggers. Expand only after the first workflows are reliable.

19. Definition of Done

BrandPilot is ready for a pilot when a team can manage approved brand information, create and review content, identify unsupported or outdated claims, complete required approvals, publish through a verified integration, handle supported engagement triggers safely, compare results with source metrics and audit who changed or approved each version.

The pilot must also show that client data is isolated, uncertainty is labelled, failed actions are visible, automation can be stopped, and time saved is measured against a documented baseline.

20. First Decisions Before Coding

- Choose the first buyer: small marketing agencies managing several brands, or an in-house marketing team. Do not build for every market at once.
- Pick one social channel and verify its exact API permissions.
- Define the first five factual/brand checks and the evidence each requires.
- Specify which decisions require a human and who can approve them.
- Set confidentiality, data-retention and AI-provider rules before using real client data.
- Run a pilot and calculate actual time saved, correction rates, automation reliability and operating cost.

Project rule: if a feature cannot be tied to a real user job, a defined control, a measurable result or a test, it does not enter the first release merely to make the product appear bigger.